

COMP3000 - Exercise 3 (Individual)

Process Manager - Scheduler - Scheduling

Winter 2018

Modify MINIX scheduling such that it implements fair-share scheduling. With fair-share scheduling, the CPU is shared by a set of processes according to the group to which they belong. For example, if there are three groups G1, G2 and G3. Each group should get one third of the CPU time with fair-share scheduling. Let us suppose that there is one process in G1, three processes in G2, and six processes in G3. The process in group G1 should get 33% of the CPU time, each process in G2 should get 11% of the CPU time and each process in G3 should get 5.5% of the CPU time.

Modify the scheduler such that all processes use fair-share scheduling. Group membership is determined by the process group identification (field `mp_procgrp` of `struct mproc`), file `/usr/src/minix/servers/pm/mproc.h`.

Use the following two test programs:

```
/* File: tester1.c
*
* A CPU-bound process (no children).
*/
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main(int argc, char *argv[]) {
    int i, j;
    i = 1;
    while (1) {
        i = j; j = i;
    }
}

/* File: tester2.c
*
* A CPU-bound process with one child.
*/
#include <stdio.h>
```

```

#include <sys/types.h>
#include <unistd.h>
int main(int argc, char *argv[]) {
    int i, j;
    fork();
    i = 1;
    while (1) {
        i = j; j = i;
    }
}

```

Start each program in different terminals using `./tester1 &` and `./tester2 &`. The command `ps` can be used to confirm that processes are created in two different groups:

```

minix# ps -o pid,ppid,pgid,command
PID PPID PGID COMMAND
667 666 667 -sh
750 667 667 ./tester2
751 750 667 ./tester2
678 677 678 -sh
749 678 678 ./tester1

```

The field `pgid` stands for process group ID.

In original MINIX, the command `top` shows that each of the CPU-bound process gets approximately one third of the CPU time.

```

load averages: 3.00, 3.00, 2.34;          up 0+01:14:50      19:59:13
52 processes: 4 runnable, 48 sleeping
CPU states: 100% user, 0.0% nice, 0.0% system, 0.0% kernel, 0.0% idle
Memory: 1023M Total, 954M Free, 953M Contig, 46M Cached

```

PID	USERNAME	PRI	NICE	SIZE	RES	STATE	TIME	WCPU	CPU	COMMAND
749	root	15	0	1224K	720K	RUN	3:23	33.79%	33.79%	tester1
751	root	14	0	1224K	720K	RUN	3:00	32.28%	32.28%	tester2
750	root	14	0	1224K	720K	RUN	2:56	30.52%	30.52%	tester2

The process associated with `tester1` is one group. When started in a different terminal, the two processes associated with `tester2` is another group. In MINIX together with your implementation of fair-share scheduling, the command `top` should show that each of the group gets approximately one half of the CPU time, while each of the two processes in the second group should get approximately of fourth of the CPU time.

```

minix# ps -o pid,ppid,pgid,command
PID PPID PGID COMMAND
687 631 631 ./tester2
688 687 631 ./tester2
686 682 682 ./tester1

```

There are two groups. Each of them should get close to 50% of the CPU time, while each of the two processes in the second group should get close to 25% of the CPU time.

```
load averages:  3.00,  2.70,  1.38;                up 0+00:51:03
54 processes: 4 runnable, 50 sleeping
CPU states: 99.7% user,  0.0% nice,  0.3% system,  0.0% kernel,  0.0% idle
Memory: 1023M Total, 124K Free, 48K Contig, 978M Cached

```

PID	USERNAME	PRI	NICE	SIZE	RES	STATE	TIME	WCPU	CPU	COMMAND
686	root	15	0	1224K	708K	RUN	2:14	47.12%	47.12%	tester1
687	root	15	0	1224K	708K	RUN	1:09	25.39%	25.39%	tester2
688	root	15	0	1224K	708K	RUN	1:08	23.97%	23.97%	tester2

Due date: February 4. This exercise must be done in the C programming language under MINIX 3.4. Submit your work on cuLearn. Submit a single *tar.gz* file. Include a README.txt file containing a report about your work (describe every change you made to the system code and where you made it). You are responsible for the completeness of your submission. You are responsible for submitting your work on time. Your submission must include a screenshot showing evidence that you code is working (see the attached example). Open three terminals, call the following sequence of commands:

1. `date;uname -v;more tester1.c;./tester1`
2. `date;uname -v;more tester2.c;./tester2`
3. `top`

```
michelbarbeau — ssh -l root -p 2222 localhost — 80x64
minix# date;uname -v;more tester1.c;./tester1
Mon Jan 22 18:24:58 GMT 2018
Minix 3.4.0 (GENERIC)
/* File: tester1.c
*
* A CPU-bound process (no children).
*/
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main(int argc, char *argv[]) {
    int i, j;
    i = 1;
    while (1) {
        i = j; j = i;
    }
}
```

```
michelbarbeau — ssh -l root -p 2222 localhost — 80x63
minix# date;uname -v;more tester2.c;./tester2
Mon Jan 22 18:24:59 GMT 2018
Minix 3.4.0 (GENERIC)
/* File: tester2.c
*
* A CPU-bound process with one child.
*/
#include <stdio.h>
#include <sys/types.h>
#include <unistd.h>
int main(int argc, char *argv[]) {
    int i, j;
    fork();
    i = 1;
    while (1) {
        i = j; j = i;
    }
}
```

michelbarbeau — ssh -l root -p 2222 localhost — 80x64

load averages: 3.00, 2.14, 1.19; up 0+00:15:48 18:27:12

54 processes: 4 runnable, 50 sleeping

CPU states: 99.7% user, 0.0% nice, 0.0% system, 0.3% kernel, 0.0% idle

Memory: 1023M Total, 968M Free, 968M Contig, 31M Cached

PID	USERNAME	PRI	NICE	SIZE	RES	STATE	TIME	WCPU	CPU	COMMAND
731	root	15	0	1224K	692K	RUN	1:06	47.82%	47.75%	tester1
735	root	15	0	1224K	692K	RUN	0:33	24.50%	24.46%	tester2
736	root	15	0	1224K	692K	RUN	0:33	24.35%	24.32%	tester2
0	root	0	0	3264K	3264K	kernel	0:00	0.05%	0.05%	[kernel]
10	root	7	0	432K	432K	RUN	0:00	0.00%	0.00%	mib
697	root	7	0	4420K	2336K	select	0:00	0.00%	0.00%	sshd
675	root	7	0	4420K	2328K	select	0:00	0.00%	0.00%	sshd
663	root	7	0	4420K	2260K	select	0:00	0.00%	0.00%	sshd
634	service	7	0	4396K	1896K	select	0:00	0.00%	0.00%	sshd
265	service	7	0	1304K	1304K	(any)	0:00	0.00%	0.00%	lwip
96	root	7	0	1168K	1168K	(any)	0:00	0.00%	0.00%	is
555	root	7	0	2232K	840K	select	0:00	0.00%	0.00%	syslogd
475	root	7	0	944K	644K	select	0:00	0.00%	0.00%	dhcpcd
664	root	7	0	884K	612K	wait	0:00	0.00%	0.00%	sh
676	root	7	0	880K	608K	wait	0:00	0.00%	0.00%	sh
698	root	7	0	876K	600K	wait	0:00	0.00%	0.00%	sh
41	root	7	0	564K	560K	(any)	0:00	0.00%	0.00%	procs
739	root	7	0	772K	528K	sysctl	0:00	0.00%	0.00%	top
1	root	7	0	428K	428K	wait	0:00	0.00%	0.00%	init
281	root	7	0	476K	288K	(any)	0:00	0.00%	0.00%	uds
114	root	7	0	280K	272K	(any)	0:00	0.00%	0.00%	devman
18	root	7	0	272K	272K	(any)	0:00	0.00%	0.00%	pci
302	root	7	0	528K	268K	pause	0:00	0.00%	0.00%	cron
661	root	7	0	420K	244K	tty	0:00	0.00%	0.00%	getty
662	root	7	0	420K	244K	tty	0:00	0.00%	0.00%	getty
659	root	7	0	420K	244K	tty	0:00	0.00%	0.00%	getty
660	root	7	0	420K	244K	tty	0:00	0.00%	0.00%	getty
643	root	7	0	544K	220K	select	0:00	0.00%	0.00%	inetd
159	root	7	0	308K	216K	select	0:00	0.00%	0.00%	devmand
285	root	7	0	208K	208K	(any)	0:00	0.00%	0.00%	ipc
33	service	7	0	204K	204K	(any)	0:00	0.00%	0.00%	at_wini
261	service	7	0	204K	204K	(any)	0:00	0.00%	0.00%	pty
256	root	7	0	168K	168K	(any)	0:00	0.00%	0.00%	lance
173	service	7	0	140K	140K	(any)	0:00	0.00%	0.00%	random
30	service	7	0	140K	140K	(any)	0:00	0.00%	0.00%	floppy
117	service	7	0	116K	116K	(any)	0:00	0.00%	0.00%	ptyfs
298	root	7	0	112K	112K	(any)	0:00	0.00%	0.00%	vbox
99	root	7	0	108K	108K	(any)	0:00	0.00%	0.00%	readclock.dr
300	root	7	0	184K	76K	select	0:00	0.00%	0.00%	update
108	service	5	0	24M	24M	(any)	0:00	0.00%	0.00%	mfs
53	service	5	0	9540K	9540K	(any)	0:00	0.00%	0.00%	mfs
7	root	5	0	1516K	1280K	(any)	0:00	0.00%	0.00%	vfs
111	service	5	0	448K	448K	(any)	0:00	0.00%	0.00%	mfs
12	service	5	0	344K	176K	(any)	0:00	0.00%	0.00%	pfs
4	root	4	0	1592K	1592K	(any)	0:00	0.00%	0.00%	rs
5	root	4	0	620K	620K	(any)	0:00	0.00%	0.00%	pm
3	root	4	0	216K	216K	(any)	0:00	0.00%	0.00%	ds
6	root	4	0	104K	104K	(any)	0:00	0.00%	0.00%	sched
8	root	3	0	132K	132K	(any)	0:00	0.00%	0.00%	memory
11	root	2	0	6072K	6072K	(any)	0:00	0.00%	0.00%	vm
289	service	2	0	164K	164K	(any)	0:00	0.00%	0.00%	log
9	root	1	0	208K	188K	(any)	0:00	0.00%	0.00%	tty
20	service	1	0	112K	112K	(any)	0:00	0.00%	0.00%	input
22	service	1	0	108K	108K	(any)	0:00	0.00%	0.00%	pckbd

